

1. is your browser running HTTP version 1.0 or 1.1? What version of HTTP is the server running?

We are using HTTP 1.1 (both source and destination)

```
▼ [Expert Info (Chat/Sequence): HTTP/
[HTTP/1.1 200 OK\r\n]
[Severity level: Chat]
[Group: Sequence]
Response Version: HTTP/1.1
```

2. What languages (if any) does your browser indicate that it can accept to the server? In the captured session, what other information (if any) does the browser provide the server with regarding the user/browser?

We are accepting Swedish and English. We are also accepting html, xhtml, images of different file formats

```
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=
Accept-Encoding: gzip, deflate\r\n
Accept-Language: sv-SE,sv;q=0.9,en-US;q=0.8,en;q=0.7\r\n
\r\n
```

3. What is the IP address of your computer? Of the gaia.cs.umass.edu server?

Our ip address is 10.241.200.65 and the destination is 128.119.245.12.

```
10.241.200.65      128.119.245.12
```

4. What is the status code returned from the server to your browser?

We got 200 OK. after the TCP packet was delivered

```
0 HTTP/1.1 200 OK (text/html
```

5. When was the HTML file that you are retrieving last modified at the server?

It is today 2023-03-27 and the it was last modified 27 mars 2023, 05:59:01 GMT

```
Server: Apache/2.4.18 (Ubuntu)
Last-Modified: Mon, 27 Mar 2023 05:59:01 GMT\r\n
ETag: "00-5f74b700dd501"\r\n
```

6. How many bytes of content are being returned to your browser?

128

```
Accept-Ranges: bytes\r\n
✓ Content-Length: 128\r\n
  [Content length: 128]
Keep-Alive: timeout=5, max=100\r\n
```

7. By inspecting the raw data in the packet content pane, do you see any HTTP headers within the data that are not displayed in the packet-listing window? If so, name one.

We got several fields that are not included into the content-pane such as content-length, the “real” host name (not the ip, the dns translated), what language we are preferring and much more

```
> GET /wireshark-labs/HTTP-wireshark-file1.html HTTP/1.1\r\n
Host: gaia.cs.umass.edu\r\n
Connection: keep-alive\r\n
Upgrade-Insecure-Requests: 1\r\n
User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q
Accept-Encoding: gzip, deflate\r\n
Accept-Language: sv-SE,sv;q=0.9\r\n
\r\n
[Full request URI: http://gaia.cs.umass.edu/wireshark-labs/HTTP-wireshark-file1.html]
[HTTP request 1/1]
```

Brief on section A 1-7:

This section is mostly about basic information about headers in requests. We learned how to find basic information about GET requests, find content-length, see when a html was updated but also what languages we prefer and also what HTTP version we use and what version the server uses.

8. Inspect the contents of the first HTTP GET request from your browser to the server. Do you see an “IF-MODIFIED-SINCE” line in the HTTP GET?

We didn't get it first time:

File data: 371 bytes  
v Line-based text data: text/html (10 lines)

```
\n
<html>\n
\n
Congratulations again! Now you've downloaded the file lab2-2.html. <br>\n
This file's last modification date will not change. <p>\n
Thus if you download this multiple times on your browser, a complete copy <
will only be sent once by the server due to the inclusion of the IN-MODIFIED
field in your browser's HTTP GET request to the server.\n
\n
</html>\n
```

but when we tried again we got the following:

```
Accept-Language: sv-SE,sv;q=0.9,en-US;q=0.8,en;q=0.7\r\n
If-None-Match: "173-5f7db708dd231"\r\n
If-Modified-Since: Mon, 27 Mar 2023 05:59:01 GMT\r\n
\r\n
```

9. Inspect the contents of the server response. Did the server explicitly return the contents of the file? How can you tell?

We would have gotten a packet with the html file but we only got a 304 returned to us instead. but the first time we can see that we get the response 200 ok and also a packet with the html file in it.

10. Now inspect the contents of the second HTTP GET request from your browser to the server. Do you see an "IF-MODIFIED-SINCE:" line in the HTTP GET? If so, what information follows the "IF-MODIFIED-SINCE:" header?

check answer 8.

11. What is the HTTP status code and phrase returned from the server in response to this second HTTP GET? Did the server explicitly return the contents of the file? Explain.

We got a 304 and we didn't get a returned html file since we already had it cached on our computer.

Brief:

Here we can see if we have retrieved a html file it stays cached on our computer until we either clean the cache or the server notifies that the html file has changed. So we can see that this saves us bandwidth. If we don't need to get the updated version we use the cached version instead.

tcp.stream eq 21						
No.	Time	Source	Destination	Protocol	Length	Info
248	29.491502	10.241.200.189	128.119.245.12	TCP	66	57379 → 80 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 WS=256 SACK_PERM
249	29.615092	128.119.245.12	10.241.200.189	TCP	66	80 → 57379 [SYN, ACK] Seq=0 Ack=1 Win=29200 Len=0 MSS=1382 SACK_PERM WS=128
250	29.615200	10.241.200.189	128.119.245.12	TCP	54	57379 → 80 [ACK] Seq=1 Ack=1 Win=131072 Len=0
251	29.615535	10.241.200.189	128.119.245.12	HTTP	526	GET /wireshark-labs/HTTP-wireshark-file3.html HTTP/1.1
252	29.735653	128.119.245.12	10.241.200.189	TCP	56	80 → 57379 [ACK] Seq=1 Ack=473 Win=30336 Len=0
253	29.735653	128.119.245.12	10.241.200.189	TCP	4200	80 → 57379 [ACK] Seq=1 Ack=473 Win=30336 Len=4146 [TCP segment of a reassembled PDU]
254	29.735653	128.119.245.12	10.241.200.189	HTTP	769	HTTP/1.1 200 OK (text/html)
255	29.735745	10.241.200.189	128.119.245.12	TCP	54	57379 → 80 [ACK] Seq=473 Ack=4862 Win=131072 Len=0
275	34.740127	128.119.245.12	10.241.200.189	TCP	56	80 → 57379 [FIN, ACK] Seq=4862 Ack=473 Win=30336 Len=0
276	34.740173	10.241.200.189	128.119.245.12	TCP	54	57379 → 80 [ACK] Seq=473 Ack=4863 Win=131072 Len=0

  

> Frame 253: 4200 bytes on wire (33600 bits), 4200 bytes captured (33600 bits) on interface \Device\NPF{...}	0020 c8 bd 00 50 e0 23 7f ac dd 50 b9 d7 cb a6 10 ...P-#- .P-...
> Ethernet II, Src: Fortinet_09:00:22 (00:09:0f:09:00:22), Dst: CloudNet_8e:6c:ad (dc:e9:94:8e:6c:ad)	0030 00 ed d0 49 00 00 48 54 54 50 2f 31 2e 31 20 32 ...I-..HT TP/1.1
> Internet Protocol Version 4, Src: 128.119.245.12, Dst: 10.241.200.189	0040 30 30 20 4f 4b 0d 0a 44 61 74 65 3a 20 57 65 64 00 OK..D ate: W
> Transmission Control Protocol, Src Port: 80, Dst Port: 57379, Seq: 1, Ack: 473, Len: 4146	0050 2c 20 32 39 20 4d 61 72 20 32 30 32 33 20 31 31 , 29 Mar 2023
Source Port: 80	0060 3a 35 37 3a 35 35 20 47 4d 54 0d 0a 53 65 72 76 :57:55 G MT- Se
Destination Port: 57379	0070 65 72 3a 20 41 70 61 63 68 65 2f 32 2e 34 2e 36 er: Apac he/2.4
[Stream index: 21]	0080 20 28 43 65 6e 74 4f 53 29 20 4f 70 65 6e 53 53 (CentOS ) Open
[Conversation completeness: Complete, WITH_DATA (31)]	0090 4c 2f 31 2e 30 2e 32 6b 2d 66 69 70 73 20 50 48 L/1.0.2k -fips
[TCP Segment Len: 4146]	00a0 50 2f 37 2e 34 2e 33 33 20 6d 6f 64 5f 70 65 72 P/7.4.33 mod_p
Sequence Number: 1 (relative sequence number)	00b0 6c 2f 32 2e 30 2e 31 31 20 50 65 72 6c 2f 76 35 1/2.0.11 Per-I/
Sequence Number (raw): 2142035280	00c0 2e 31 36 2e 33 0d 0a 4c 61 73 74 2d 4d 6f 64 69 -16.3-1 ast-Mo
[Next Sequence Number: 4147 (relative sequence number)]	00d0 66 69 65 64 3a 20 57 65 64 2c 20 32 39 20 4d 61 fied: We d, 29
Acknowledgment Number: 473 (relative ack number)	00e0 72 20 32 30 32 33 20 30 35 3a 35 39 3a 30 32 20 r 2023 0 5:59:0
Acknowledgment number (raw): 3117927334	00f0 47 4d 54 0d 0a 45 54 61 67 3a 20 22 31 31 39 34 GMT-ETa g: "11
0101 .... = Header Length: 20 bytes (5)	0100 2d 35 66 38 30 33 61 63 34 38 66 33 36 34 22 0d -5f803ac 48f364
> Flags: 0x010 (ACK)	0110 0a 41 63 63 65 70 74 2d 52 61 6e 67 65 73 3a 20 -Accept- Ranges
Window: 237	0120 62 79 74 65 73 0d 0a 43 6f 6e 74 65 6e 74 2d 4c bytes..C ontent
[Calculated window size: 30336]	0130 65 6e 67 74 68 3a 20 34 35 30 30 0d 0a 4b 65 65 length: 4 500- K
[Window size scaling factor: 128]	0140 70 2d 41 6c 69 76 65 3a 20 74 69 6d 65 6f 75 74 p-Alive: timeo
	0150 3d 35 2c 20 6d 61 78 3d 31 30 30 0d 0a 43 6f 6e =5, max= 100- C
	0160 6e 65 63 74 69 6f 6e 3a 20 4b 65 65 70 2d 41 6c nection: Keep-
	0170 69 76 65 0d 0a 43 6f 6e 74 65 6e 74 2d 54 79 70 ive- Con tent-T

12. How many HTTP GET request messages did your browser send? Which packet number in the trace contains the GET message for the Bill or Rights?

We sent one get request, the package number is 109

13. Which packet number in the trace contains the status code and phrase associated with the response to the HTTP GET request? What is the status code and phrase in the response?

114 contains the 200 ok.

14. How many data-containing TCP segments were needed to carry the single HTTP response and the text of the Bill of Rights?

we needed two packages, the first one contained 4200 bytes and the seconds contained 769 bytes

```

  [2 Reassembled TCP Segments (4861 bytes): #113(4146), #114(715)]
    [Frame: 113, payload: 0-4145 (4146 bytes)]
    [Frame: 114, payload: 4146-4860 (715 bytes)]
    [Segment count: 2]
    [Reassembled TCP length: 4861]
    [Reassembled TCP Data: 485454502f312e3120323030204f4b0d0a446174653a204d6f6e2c203237204d61722032...]
  [Hypertext Transfer Protocol]

```

15. How many data-containing TCP segments were needed to carry the single HTTP response and the text of the Bill of Rights?

No, the package is not segmented hence they have a big packet with the full html file with len=4146. We can see the rest of the files have len=0 so they come with just a header with 0 content. Check the first image of this section.

Brief: Here we can see the difference between pipelined requests and serially send get requests. We can also conclude that we can save a lot of time using pipelines if we are trying to get multiple pictures from multiple sites. We can also see, if a package is too large we can see the package getting split up into several packages with the same identifier.

16. How many HTTP GET request messages were sent by your browser? To which Internet addresses were these GET requests sent?

We are doing three GET requests, first we retrieve the html then we try to retrieve two images. The HTML and the first 200 OK image are retrieved from the website we entered. The second picture is retrieved from a new ip address ("178.79.137.164"). We can see the ip address corresponds to kurose.cslash.net with the link /8E\_cover\_small.jpg there for the link must be: [kurose.cslash.net/8E\\_cover\\_small.jpg](http://kurose.cslash.net/8E_cover_small.jpg)

No.	Time	Source	Destination	Protocol	Length	Info
86	6.547870	10.241.200.65	128.119.245.12	HTTP	526	GET /wireshark-labs/HTTP-wireshark-file4.html HTTP/1.1
90	6.671647	128.119.245.12	10.241.200.65	HTTP	1355	HTTP/1.1 200 OK (text/html)
91	6.692662	10.241.200.65	128.119.245.12	HTTP	472	GET /pearson.png HTTP/1.1
96	6.815639	128.119.245.12	10.241.200.65	HTTP	901	HTTP/1.1 200 OK (PNG)
110	7.399912	10.241.200.65	178.79.137.164	HTTP	439	GET /8E_cover_small.jpg HTTP/1.1
112	7.431365	178.79.137.164	10.241.200.65	HTTP	225	HTTP/1.1 301 Moved Permanently

```
> GET /8E_cover_small.jpg HTTP/1.1\r\n
Host: kurose.cslash.net\r\n
Connection: keep-alive\r\n
User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64
```

17. Can you tell whether your browser downloaded the two images serially, or whether they were downloaded from the two websites in parallel? Explain.

We can see this was retrieved serially since we first get a 200ok for the html, then we make a new request. We get again 200 Ok then we do the last Get request and get the 200 ok. If this was made in parallel we would do two get requests at the same time. We can also see on the time elapsed that this can't be made in parallel since there is around 0,4 seconds between the requests which is a lot of time keeping in mind we are talking about computers.

Brief: Here we can see that we can retrieve different files from different sites. Everything doesn't need to be on the same site. This is why proxies exist and can save us loads of time if we need to travel for a long distance. (getting an image from a proxy in Europe is way faster than trying to get it from the original server in Canada).

No.	Time	Source	Destination	Protocol	Length	Info
73	6.993473	10.241.200.65	128.119.245.12	HTTP	542	GET /wireshark-labs/protected_pages/HTTP-wireshark-file5.html HTTP/1.1
75	7.114971	128.119.245.12	10.241.200.65	HTTP	771	HTTP/1.1 401 Unauthorized (text/html)
208	18.858198	10.241.200.65	128.119.245.12	HTTP	627	GET /wireshark-labs/protected_pages/HTTP-wireshark-file5.html HTTP/1.1
210	18.981165	128.119.245.12	10.241.200.65	HTTP	544	HTTP/1.1 200 OK (text/html)

18. What is the server's response (status code and phrase) in response to the initial HTTP GET message from your browser?

We get a 401 unauthorized since we haven't logged in.

19. When your browser sends the HTTP GET message for the second time, what new field is included in the HTTP GET message?

The new fields is the login form we now retrieve since the server knows that we are not logged in.

Brief: Here we can see the protected site returning a 401 unauthorized firstly since the server assumes we are logged in. When we receive the 401 we send a new package asking for the login form, which we get and after that we can send our information and later get a 200 OK.

20. What does the "Connection: close" and "Connection: keep-alive" header field imply in HTTP protocol? When should one be used over the other?

Connection: close means the client wants the server to close the connection after sending the response, while Connection: keep-alive means the client wants the connection to remain open for multiple requests. Choosing between the two depends on the specific use case. Connection: close is more efficient for a single request, while Connection: keep-alive is better for multiple requests over the same connection.